

Symbolic Model Checking of Timed Automata using LTSmin

Sybe van Hijum

August 3, 2016

Contents

1	Introduction	3
2	Preliminaries	4
2.1	Timed Automata	4
2.2	Zones	5
2.3	Zone subsumption	5
3	Related Work	7
3.1	Methods	7
3.2	LTSmin	11
3.3	Difference Decision Diagrams	11
3.4	List Decision Diagram	17
4	Implementation	18
4.1	Flattening DBM	18
4.2	Dependency Matrices	18
4.3	DBM reduction	19
4.4	Connecting LDD and DDD	19
4.5	DDD nodes	21
4.6	Creating Nodes	21
4.7	Minus	22
5	Benchmarks	26
5.1	Viking	26
5.2	Fischer	26
5.3	CSMA-CD	26
5.4	Animo	27
5.5	Lynch-Shavit	27
5.6	Milner	27
5.7	Other models	27
6	Results	28

1 Introduction

Timed automata [2] is a widely used modelling formalism. A recent usage of this formalism is the modelling of biological signalling pathways [27]. ANIMO is a tool that generates these timed automata from biological signalling pathways models. This leads however to large state spaces, and sometimes to models that are too large to handle by conventional methods. Therefore better model checking techniques for timed automata, that can handle larger state spaces are needed. We look into symbolic algorithms for timed automata.

BDDs (Binary Decision Diagrams) [1, 10] and variations like LDDs (List Decision Diagrams) [9] and MDDs (Multi-valued Decision Diagrams) [28] have proven their value in model checking algorithms. Due to advances in this field, models with much larger state spaces can be explored on the same machine. This progress has not been translated directly to more efficient methods for timed automata. Several methods have been proposed, like CDDs (Clock Difference Diagrams) [19], CMDs (Constraint Matrix Diagrams) [15], CRDs (Clock Restriction Diagrams) [30] and DDDs (Difference Decision Diagrams) [23, 25]. All of these methods show some extra difficulties or limitations over BDDs. Also after their introduction they have not been developed further.

LTSmin [8, 17] is a language independent on the fly model checker with several algorithmic back-ends. Its symbolic back-end uses BDDs to both represent the state space and the transition relations of models. These BDDs are generated on the fly by the search algorithms. LTSmin has a language module for Uppaal [4] through the Opaal [12] lattice model checker. Through this module Uppaal models can be loaded into LTSmin. For this language currently, only the multi-core back-end can be used [11]. This multi-core approach showed efficient enough to compete with the latest version of the Uppaal model checker. It showed significant speedups on multi-core machines, at the cost of some memory increase however. To tackle the memory increase a combination of the Opaal front-end and the symbolic back-end could be a solution.

The symbolic back-end of LTSmin provides both a memory reduction by using BDDs and a speedup by using multi-threaded search algorithms and the multi-threaded BDD package Sylvan [29]. Using this together with the Uppaal language front-end will hopefully result in a model checker that can compete both on time and memory consumption with the Uppaal model checker.

We will propose a symbolic reachability for timed automata that is capable of handling the models that are generated by the ANIMO tool.

2 Preliminaries

We will first define timed automata and zones, a method used to represent time in timed automata. Also a subsumption check over zones will be defined.

2.1 Timed Automata

Timed automata is a formalism that extends labelled transition systems with one or more clocks. Guards over these clocks, denoted as $G(C)$ can be used for transitions. Also reset actions for clock can be defined for transitions. All clocks in the system will increase at the same rate. As our work continues on [11] we use the same definition of timed automata.

Definition 1 (Timed Automata). *An extended timed automaton is a 6-tuple $A = \langle L, C, Act, s_0, \rightarrow, I_C \rangle$ where*

- L is a finite set of locations, typically denoted by l
- C is a finite set of clocks, typically denoted by c
- Act is a finite set of actions
- $l_0 \in L$ is the initial location
- $\rightarrow \subseteq L \times G(C) \times Act \times 2^C \times L$ is the (non-deterministic) transition relation. We normally write $l \xrightarrow{g,a,r} l'$ for a transition., where l is the source location, g is the guard over the clocks, a is the action, and r is the set of clocks reset.
- $I_C : L \rightarrow G(C)$ is a function mapping locations to downwards closed clock invariants.

With this definition we can combine timed automata to a network of timed automata, which is a parallel composition, to define larger systems.

Definition 2 (Network of timed automata [11]). *Let $Act = \{ch!, ch? | ch \in Chan\} \cup \{\tau\}$ be a finite set of actions, and let C be a finite set of clocks. Then the parallel composition of extended timed automata $A_i = \langle L_i, C, Act, S_0^i, \rightarrow_i, I_C^i \rangle$ for all $1 \leq i \leq n$, where $n \in \mathbb{N}$, is a network of timed automata, denoted $A = A_1 || A_2 || \dots || A_n$.*

A network of timed automata is a parallel composition that synchronizes on a set of channels $Chan$ [4]. $ch!$ and $ch?$ represent the output and input action on the channel $ch \in Chan$.

$$\begin{array}{c}
\mathbf{O} \qquad c_1 \qquad c_2 \\
\mathbf{O} \begin{pmatrix} (0, \leq) & (0, \leq) & (0, \leq) \\ c_1 \begin{pmatrix} (5, \leq) & (0, \leq) & (\infty, \leq) \\ c_2 \begin{pmatrix} (4, \leq) & (\infty, \leq) & (0, \leq) \end{pmatrix} \end{pmatrix} \end{pmatrix}
\end{array}$$

Figure 1: DBM

2.2 Zones

For basic transition systems the state space can grow exponentially for the size of the system. The state space of Timed automata is by definition infinite, as clocks have real values. If a state is defined between two points in time, an infinite amount of moments in time can happen during that state. Even when some granularity is used, that defines that clocks will only increase with certain step size the automata can still have infinite state space if a clock is unbounded. To tackle this problem most model checkers use a notion of zones for the representation of time. A zone can be seen as a set of constraints over the clocks C of the form $c_i \sim x$ and $c_i - c_j \sim x$ where $\sim \in \{<, \leq, =, \geq, >\}$ and $x \in \mathbb{N}$. To represent these zones several data structures have been developed. One of the most common used structures are Difference Bound Matrices (DBMs) [5, 13].

These matrices use both a column and a row for each clock, and on each position (i, j) an upper bound on the difference between the clocks c_i and c_j is given in the form $c_i - c_j \preceq x$ where $\preceq \in \{<, \leq\}$ and $x \in \mathbb{Z}$. For the constraints over the single clocks an extra clock \mathbf{O} with a constant value 0 is added. This way the upper and lower bound of a clock c_i can be given by $c_i - \mathbf{O} \preceq x$ and $\mathbf{O} - c_i \preceq y$. The addition of this \mathbf{O} clock will give the matrix of a timed automaton always size $(|C| + 1)^2$. This way convex zones of clock variables can be represented. Each matrix can however only contain a single convex zone. Concave zones and multiple convex zones need multiple matrices to be represented. As a solution often a list of DBMs is used. In figure 1 we give an example of a DBM with two clocks: c_1 and c_2 , representing the zone $0 \leq c_1 \leq 5 \wedge 0 \leq c_2 \leq 4$. The diagonal only contains $(0, \leq)$ values as these elements give the difference between a clock and itself, which is clearly always 0.

2.3 Zone subsumption

In model checking an important function is to check if a certain state has been visited already earlier. For normal automata this can be done by comparing the newly found state to all states that have already been visited, and check if one of those states is equal to that new state. This is often done by more efficient methods, like hash functions, but the equality check remains. For states with zones this equality check does not satisfy. Two

zones do not need to be equal, but the newly discovered zone can also be a subset of the earlier discovered zones. In LTSmin this is done by a subsumption check [11] that is performed over the DBMs. This check is delegated to the Uppaal DBM library. The function checks if a new zone is a subset of the zone represented by a DBM.

3 Related Work

In this related work section we will discuss a number of methods used for model checking timed automata. We will choose a method to extend our work on, and go more into detail on that method.

3.1 Methods

Already several model checkers for timed automata exist such as Uppaal [4], KRONOS [31], RABBIT [7] and RED [30]. We focus mainly on the Uppaal tool as we use the same input format. Opaal [12], the language module for LTSmin, uses the XML format that is created by the Uppaal tools. This way we can use the Uppaal user interface to create and adapt models. We also use the Uppaal DBM library to represent zones. Several methods exist to represent the clock variables in a timed model. The most used methods are digitization and zones.

Digitization approximates the continuous values of clocks by using discrete values [6]. The method however only works for closed timed automata, meaning that no strict comparisons on clocks can be made in the model and that clocks only can be compared to integers. This approach is very sensitive to the granularity of the values used and the upper bound of the clock values. When fine granularity or large upper bounds are used, the memory usage will increase too much. An advantage of this approach is that basic model checking approaches can be used and no extra complexity due to zone calculations is added. This method results in a transition system with only discrete variables, so a normal BDD package can be used. In [26] a similar approach is proposed by using clock tick actions to represent time progress and removing clock variables altogether.

The most established method to represent clock zones are DBMs. We gave an introduction to this structure in the preliminaries section. Several methods based on BDDs have been developed to represent zones. All of these are similar to DBMs in the sense that they use clock constraints to represent the zones. These structures use a BDD-like structure to represent the zones more efficiently. Below we shortly describe four methods. For each method we give an example, all examples represent $2 < c_1 - c_2 < 4 \vee 7 \leq c_1 - c_2 \leq 8$.

CDDs [19] use single nodes for each variable and have multiple disjoint intervals for that variable on the edges. This results in a node with a larger fanout. The upper and lower bound for each pair of clocks are represented in a single node, as the edges represent intervals. Requiring the disjointness of intervals can lead to a memory inefficient representation, as intervals need to be cut in more smaller parts. All algorithms on CDDs do not maintain disjointness, after every step it needs to be re-established. In figure 2 we have an example of a CDD.

DDs [23,25] use an upper-bound constraint on each node that can either




Figure 2: CDD representation

be true or false. Each node thus has a fixed fanout of two. When a constraint is false, a next node will have another constraint on the same variable. This requires a fixed ordering based on the variables, values and operators. In figure 3 an example of a DDD is shown.




Figure 3: DDD representation

CRDs [30] differ mainly from CDDs by not using disjoint intervals but possibly overlapping upper bounds, for a pair of variables on their edges. This diagram will have a larger fanout per node, like CDDs. Several normal forms for this diagram are proposed, with different performance results. It is also shown that CRDs can be combined with BDDs into a single structure to fully symbolic represent state space. In figure 4 we give an example of a CRD.

CMDs [15] combine CDDs, CRDs and DBMs into a single structure. This diagram type differs from the others by having multiple constraints per edge, resulting in a diagram with few nodes. Upper- and lower-bounds of multiple clock pairs can be on a single edge. CMDs do not have a canonical form so only some reductions are proposed. An example of a CMD is given in figure 5. This figure contains two examples, the first is a diagram of the constraint we use in this section. To show the difference with other diagrams we also give a diagram representing the same zone as the DBM in figure 1.




Figure 4: CRD representation




Figure 5: CMD representation

In [14, 32] a method is proposed purely based on BDDs by translating the constraints directly into BDD nodes. We call this method BDD zones. This results in a unified structure for both the discrete variables and the clock constraints. The method is only a proof of concept and has not been implemented in a model checker and no performance results are known. Subsumption for this method may be difficult. On BDDs only equalities can be checked, and no inequalities. This way inclusion is not trivial to check by normal BDD algorithms.

A known difficulty in BDDs is the variable ordering. A bad ordering can lead to a BDD of exponential size, where a good ordering can sometimes lead to a significantly smaller diagram. Of the zone diagrams named above, only for CRDs experiments with different orderings have been conducted, the other researches assume a given ordering on the variables and the ordering of the values is fixed. The CRD case shows that full interleaving and having related variables close to each other in the ordering is preferable and gives the best results, both on speed and memory. This is the same result as expected with BDDs, this suggests that similar orderings should be used with these techniques. In Table 1 we compare the different types of diagrams we discussed above.

Table 1: Comparing Diagrams

Type	Pro	Con
DBM	Canonical form for convex zones Existing library Inclusion check	Concave zones need multiple DBMs Not memory efficient
DDD	Structure like LDD Re-ordering of variables possible Apply same efficiency as BDDs Boolean variables also in DDD	Canonicity hard to obtain No on the fly canonicity Expensive normal form computation Only time performance tested Only reduction algorithms
CDD	Structure like MDD Inclusion check (intersection of complement)	No algorithm to get normal form Only high level algorithms given Methods don't maintain disjointness Expensive normal form computation No implementation results available Disjointness memory inefficient
CRD	Combination with BDD possible Variable reordering shows advantage Library available Some benchmarks exp better than CDD Extensive benchmarks Good performance backwards reach	3 possible canonical forms No algorithms in paper Some benchmarks linear worse than CDD
CMD	Benchmarks against RED and Uppaal	Results differ per case Needs translation from vector to edges Two reduced forms
BDD discrete	Using existing BDD packages Good performance for small clock values	Performance decreases fast for large values Not possible with current Opaal PINS Introducing additional 'tick' actions Only for closed timed automata
BDD zones	Using existing BDD packages All variable reorderings possible Only need direct translation DBM to state vector Easy to implement	Losing zone containment No implementation results




Figure 6: Modular structure of LTSmin

3.2 LTSmin

LTSmin [8, 17] is a language independent model checker. It is built in a modular way such that new languages can be added by a PINS (Partitioned Next-State Interface) interface without too much effort, and new algorithms can be added easily. LTSmin offers four different algorithmic back-ends for model analysis: symbolic, multi-core, sequential and distributed. All of these back-ends support different types of reduction and model checking. Several language modules have already been built for LTSmin such as mCRL2, Promela, DVE and Uppaal. The modular structure of LTSmin is shown in Figure 6. The PINS is the core of LTSmin. This interface abstracts as much as possible from the model without losing the structure. It represents states as fixed length integer arrays. The main function of the interface is a (partitioned) next-state function which returns the successor states. With these functions a state space can be generated on the fly. With the use of dependency matrices event locality can be determined statically [21]. With these matrices, more efficient symbolic algorithms can be used, the number of next-state calls can be reduced, efficient variable re-orderings can be used, and transition caching can be used. In the current Uppaal PINS the next-state function is not partitioned and therefore no meaningful dependency matrix is created, and none of these algorithms can be used. Also the DBM variable is only represented by a pointer, which is not a meaningful value for the transition system. LTSmin uses the pointer to a DBM to do the subsumption check as described in section 2.3.

3.3 Difference Decision Diagrams

We have discussed several symbolic approaches for representing zones. All of these approaches have benefits and downsides over each other. We chose

to develop one of these approaches in LTSmin. We wanted a diagram that can store both discrete states and zones, this can either be done in the diagram, or in a combination of the diagram and BDD or LDD nodes. Also a subsumption check on the diagram should be possible. We chose from the four zone representing diagrams discussed earlier. The CDD approach was not chosen due to the memory inefficient disjoint intervals and their algorithms not maintaining these disjointness. The CMD approach is too similar to DBMs, on which we already have an approach. The choice between CRD and DDD was between two quite similar diagrams. We have decided to continue on the DDD. It is a diagram form that is closely related to LDDs, for which we already have a library, so we can reuse parts of the LDD library, and it is also quite compatible to the current PINS structure and its next-state function. The method still has some loose ends that need research, mostly on the algorithms and efficiently creating a canonical form. No results on the memory usage are available, which is normally the greatest benefit of a symbolic approach, so also on the results side we can extend the current research.

So DDDs are a diagram type that seems to fit well in the current structure we have, but there is still room for some more research. First we give the definition of a DDD.

Definition 3 (Difference Decision Diagram [25]). *A difference decision diagram (DDD) is a directed acyclic graph (V, E) . The vertex set V contains two terminals 0 and 1 with out-degree zero, and a set of non-terminal vertices with out-degree two and the following attributes.*

Attribute	Type	Description
$pos(v), neg(v)$	Var	Positive variable x_i , and negative variable x_j .
$op(v)$	$\{<, \leq\}$	Operator $<$ or \leq .
$const(v)$	\mathbb{D}	Constant c .
$high(v), low(v)$	V	High-branch h , and low-branch l .

The set E contains the edges $(v, low(v))$ and $(v, high(v))$, where $v \in V$ is a non-terminal vertex.

Now we have the definition of the structure. We also give the semantics of this structure.

Definition 4 (DDD semantics). *The semantics of a vertex is defined recursively by the function $\mathcal{V} : V \rightarrow \mathbf{Exp}$:*

- $\mathcal{V}[[0]] \stackrel{\text{def}}{=} \text{false},$
- $\mathcal{V}[[1]] \stackrel{\text{def}}{=} \text{true},$

$$\bullet \mathcal{V}[[v]] \stackrel{\text{def}}{=} \begin{cases} (pos(v) - neg(v) < const(v)) \rightarrow \mathcal{V}[[high(v)]], \mathcal{V}[[low(v)]] \text{ if } op(v) = '<' \\ (pos(v) - neg(v) \leq const(v)) \rightarrow \mathcal{V}[[high(v)]], \mathcal{V}[[low(v)]] \text{ if } op(v) = '\leq' \end{cases}$$

In the semantics we only take the information on the high edges. The implicit information on the low edge is not used. A node can thus only represent an upper-bound which is either true or false, it can not implicitly represent a lower-bound on the same variable pair. This representation also makes it easier to work with the state-vectors of LTSmin.

In [25] a canonical form for DDDs is discussed, also called a fully reduced DDD. Only definitions are given here, no algorithms to reach this form. It is stated that it is difficult to reach this fully reduced form. It is not clear if they managed to make their apply function in such a way that it maintains canonicity, as the function for BDDs does. To reach canonicity, local reductions and ordering are a first step, but it is not enough due to dependencies among the constraints. For BDDs the local reductions and ordering are sufficient to reach a canonical form. First we give some notational shorthands and then we define an ordering and local reductions on DDDs.

$$\begin{aligned} var(v) &= (pos(v), neg(v)) \\ bound(v) &= (const(v), op(v)) \\ cstr(v) &= (var(v), bound(v)) \end{aligned}$$

To order DDD nodes we use the operator \prec . This orders variables and variable pairs in a predefined order. It orders bounds by increasing constants, and the $<$ operator before the \leq operator. So a node v with $bound(v) = (0, <)$ comes before $bound(u) = (0, \leq)$ which comes before $bound(w) = (1, <)$.

Definition 5 (Ordered DDD [25]). *An ordered DDD (ODDD) is a DDD where each non-terminal vertex v satisfies:*

1. $neg(v) \prec pos(v)$,
2. $var(v) \prec var(high(v))$,
3. $var(v) \prec var(low(v))$ or $var(v) = var(low(v))$ and $bound(v) \prec bound(low(v))$.

After ordering a DDD some local reductions can be defined to reduce the size of a DDD.

Definition 6 (Locally Reduced DDD [25]). *A locally reduced DDD (R_L DDD) is an ODDD satisfying, for all non-terminals u and v :*

1. $\mathbb{D} = \mathbb{Z}$ implies $\forall v. op(v) = '\leq'$,
2. $(cstr(u), high(u), low(u)) = (cstr(v), high(v), low(v))$ implies $u = v$,

3. $low(v) \neq high(v)$,
4. $var(v) = var(low(v))$ implies $high(v) \neq high(low(v))$.

We give an example of the last point in figure 7. Here both diagrams represent the same zone: $2 < c_1 - c_2 \leq 8$. The node with < 4 on the high edge is redundant in this example and can thus be removed.




Figure 7: Local reduction

For BDDs these reductions would be enough to have a fully canonical structure. For DDDs this is not the case. Due to dependencies between the bounds. In figure 8 we give an example for this by giving two different locally reduced DDDs representing the same zone. The resulting zone of both these DDDs is drawn in figure 9 we show the result of this zone, which is the square in which both clock c_1 and c_2 are between 0 and 5.

The R_LDDD is clearly not canonical. We first define a path in a DDD as the bound on all high edges that are traversed in a single walk from the top node to the true node. A path will only have one bound for each variable pair.

Definition 7 (Path-reduced DDD [25]). *A path-reduced DDD (R_PDDD) is a locally reduced DDD where all paths are feasible.*

This definition ensures that all paths in a DDD actually represent a zone, and that there are no redundant paths in the DDD that just represent an empty set. This usage of paths is compatible to the state vectors used in LTSmin. An R_PDDD is still not canonical. We need to define tightness, saturation and disjunctive vertices. To define tightness we first need to define dominating constraints.

Definition 8 (Dominating constraint [25]). *A constraint $x_i - x_j \lesssim c$ is dominating in a path p if all other constraints $x_i - x_j \lesssim' c'$ on the same pair of variables in p are less restrictive.*




Figure 8: Two DDDs representing the same zone

Definition 9 (Tightness [25]). A dominating constraint $\alpha = x_i - x_j \lesssim c$ is tight in a feasible path $[p] = [p_1] \wedge \alpha \wedge [p_2]$ if for all tighter constraints $(c', \lesssim') < (c, \lesssim)$, the systems $[p_1] \wedge (x_i - x_j \lesssim' c') \wedge [p_2]$ and $[p]$ have different solutions. A path p is tight if it is feasible and all dominating constraints on it are tight. An $R_L DDD u$ is tight if all paths from u are tight.

Definition 10 (Saturation [25]). A tight path p from an $R_P DDD$ is saturated if for all constraints α not on p , if α is added to p either (1) α is not dominating and tight, or (2) the constraint system $[p_1] \wedge \neg \alpha$ is infeasible when $[p]$ is written $[p] = [p_1] \wedge [p_2]$ with all constraints on p_1 smaller than α with respect to $<$ and all constraints on p_2 larger than α . An $R_P DDD u$ is saturated if all paths from u are saturated.

Definition 11 (Disjunctive vertex [25]). Let p be a path leading to the vertex u in a DDD, and assume $\alpha = cstr(u)$, $h = high(u)$, and $l = low(u)$. Then u is disjunctive in p if $[p] \wedge (\alpha \rightarrow h, l)$ and $[p] \wedge (h \vee l)$ have the same set of solutions.




Figure 9: Resulting zone of DDDs in figure 8

All of these definitions together lead to the following definition of a fully reduced DDD.

Definition 12 (Fully reduced DDD [25]). *An R_pDDD u is a fully-reduced DDD (R_FDDD) if it is tight, saturated and has no disjunctive vertices.*

We assume that this fully-reduced DDD is canonical and work from that. It is not ensured that this is actually the case, there is no proof for it.

Conjecture 1 (Canonical DDD [25]). *If u and v are R_FDDD s with the same set of solutions then $u = v$.*

DDD's can also be used to represent the discrete variables in automata. This is done by translating the variable into a difference constraint. For example $x_1 = 3$ will be translated into $x_1 - 0 \leq 3 \wedge 0 - x_1 \leq -3$, thus resulting into a DDD with two nodes. We will connect the DDD to an LDD to represent discrete variables to limit the number of nodes.

So far we only found the results of two benchmark tests of DDDs, Milner's scheduler and Fischer's protocol [24]. Here the DDD approach has been compared with KRONOS and Uppaal which were both slower than the DDD implementation. The results of these benchmarks show no memory usage or number of nodes needed.

3.4 List Decision Diagram

The DDD nodes are connected to LDD nodes to represent the discrete variables. We will introduce the LDD structure here. An LDD is used to represent variables with integer values, not only binary values. In contrast to MDDs this is done for one value per node. Resulting in nodes with equal size. We will first define the LDD structure.

Definition 13 (List Decision Diagram). *A list decision diagram (LDD) is a directed acyclic graph (V, E) . The vertex set V contains two terminals 0 and 1 with out-degree zero, and a set of non-terminal vertices with out-degree two and the following attributes.*

<i>Attribute</i>	<i>Type</i>	<i>Description</i>
$var(v)$	Var	Variable x
$const(v)$	\mathbb{Z}	Constant c .
$high(v), low(v)$	V	High-branch h , and low-branch l .

The set E contains the edges $(v, low(v))$ and $(v, high(v))$, where $v \in V$ is a non-terminal vertex.

The definition is almost equal to DDDs, definition 3. The difference is the operator that is not in LDDs. LDDs can be seen as a DDD with not a $<$ or \leq as operator, but a $=$.

4 Implementation

This section will go more into detail about the implementation we made and the design choices that were needed.

4.1 Flattening DBM

In the LTSmin implementation that we already have the state vector exists of all discrete variables and an 64 bit pointer to a C++ class containing a DBM [11]. For a symbolic solution this pointer has no meaning, thus we take the actual values from the DBM and put these into the state vector. This increases the length of a state vector, but does not need to increase the memory footprint, as the DBM was already stored. In the DBM library we use a DBM is represented by a one dimensional array of 32-bit integers. In the integers the complete bound is stored, so both the operator and the constant value. We flattened the DBMs to work with a symbolic solution. We only did this on the edges of the successor function. So this function reads a flattened DBM as input and returns it as successor states, internally the original DBM representation is still used. This way the code had to be adapted the least. In this flattening we removed the diagonal elements of each DBM. By the way DBMs are constructed this will always represent the difference between a clock and itself. This difference is by definition always 0, so it can be removed, and hard coded be set to $(0, \leq)$ internally. This reduces the number of state variables in the state vector by one for each clock. This flattening of DBMs results into a language module that can be connected to all LTSmin algorithmic back-ends for state-space generation.

4.2 Dependency Matrices

To get the best possible result of the regrouping algorithms, the dependency matrices had to be made as sparse as possible. This has been done for both the read matrix and may-write matrix. For even better results, also the must-write matrix is needed. This needs effort when analysing the code, this can be done, but is left out for this thesis. First of all, all C-like code is parsed. Here it is stored per function which variables are read and written, and which other functions are called. Next all transitions are parsed, here some variables are read and written directly. Transitions can also call functions, in such cases the variables that were found in the parsing of these functions are added to the read and may-write variables of the transition. In the third step we need to look at the time extrapolation. This extrapolation is based on the value of the location variable, so it results in a read dependency. In some cases, there is no difference between all possible location values for this extrapolation, so a location does not need to be read. A final step is that a location variable that can be urgent or committed always has

to be read. If this location is in an urgent state, than no other transitions can happen, so all other transitions have to check that they are not in an urgent state. In which only an other transition can take place. The flattened DBMs and the sparser dependency matrices together enable the reordering algorithms in the symbolic back-end of LTSmin to be used.

4.3 DBM reduction

We work towards a fully reduced DDD solution. This is already started at the language module size. The next-state function will only return tight and saturated paths. In DBM terms this is a minimal constraint system [5]. As the length of a state-vector can not be changed on the fly, all removed constraints are set to $(<, \infty)$. This means that there is no upper-bound on the variable pair of that position. In algorithm 2 which uses algorithm ?? we show the algorithm that determines all bounds that are not needed and can be set to $(<, \infty)$. The DBM library can not use these minimal constraint systems. In the next state function the incoming DBM is tightened, then all needed operations for the successor generation are conducted and if a successor is returned, its DBM is again turned into a minimal constraint system. This will give algorithmic overhead for each next-state call. The advantage of this procedure is that many bounds will be redundant and turned into $(<, \infty)$. In the symbolic back-end these bounds which are the same can be shared in a single node. Thus taking more time in the successor generator, it can also reduce the number of nodes in the algorithmic back-end. This reduction is used in the successor generator for the symbolic back-end, and will also be used for the DDD solution.

Algorithm 1 Reduce

```

1: procedure REDUCE( $dbm, dim$ )
2:   for  $i \in dim$  do
3:     for  $j \in dim$  do
4:       for  $k \in dim$  do
5:         if  $!(dbm[i, k] \vee dbm[k, j] \vee dbm[i, j] \text{ on diagonal})$  then
6:           if  $dbm[i, k] + dbm[k, j] \leq dbm[i, j]$  then
7:              $dbm[i, j] := \infty$ 

```

4.4 Connecting LDD and DDD

To represent the discrete variables in states LDD nodes are used. The structure of these nodes is quite similar to DDD nodes. We decided to not mix the nodes, but to first have all the LDD nodes and then all DDD nodes in the tree. In the state vector the first part exists of all discrete variables, the last part are the DBM variables. The top of the diagram

Algorithm 2 Reduce

```
1: procedure REDUCEZERO( $dbm, dim$ )
2:    $placed[dim]$  all 0
3:    $red[dim, dim]$  all 0
4:    $eq[dim, dim]$  all 0
5:    $cl := 0$ 
6:    $newDBM[dim, dim]$  diagonal  $\infty$  rest 0
7:   for  $i \in dim$  do
8:     if  $placed[i] = 0$  then
9:       for  $j \in dim$  do
10:        if  $dbm[i, j] + dbm[j, i] = 0$  then
11:           $placed[j] := 1$ 
12:           $eq[cl, j] := 1$ 
13:         $cl ++$ 
14:    $repr[cl]$ 
15:   for  $i \in cl$  do
16:     for  $j \in dim$  do
17:       if  $eq[i, j] = 1$  then
18:          $repr[i] := j$ 
19:       break
20:    $clg[cl, cl]$ 
21:   for  $i \in cl$  do
22:     for  $j \in cl$  do
23:        $clg[i, j] := dbm[repr[i], repr[j]]$ 
24:   REDUCE( $clg, cl$ )
25:   for  $i \in cl$  do
26:     for  $j \in dim$  do
27:       if  $eq[i, j] = 1$  then
28:         for  $k \in dim$  do
29:           if  $eq[i, k]$  then
30:              $newDBM[j, k] = dbm[j, k]$ 
31:       for  $j \in cl$  do
32:          $newDBM[repr[i], repr[j]] := clg[i, j]$ 
33:   return newDBM
```

can be seen as a MTLDD(Multi-Terminal List Decision Diagram) with not values on the leaf nodes, but pointers to DDD nodes. The DDD part is not influenced by the LDD part, as a node is only influenced by the nodes below it, it has no information about the nodes above it in the diagram. This strict separation between LDD and DDD nodes makes that the reordering algorithms can not be used. The lack of reordering makes it also possible to reconstruct the DBMs on the DDD side. This is used for the minus function which we discuss later.

4.5 DDD nodes

We used the basis of the LDD package in Sylvan to create our DDD nodes. The nodes are the same as the LDD nodes, only two previously unused bits are now used to store the operator and the type of the node. DDD nodes are stored in 128 bits, represented as a struct of two 64 bit integers. The hashtable that is already used by Sylvan is specifically for 128 bit entries, so the DDD nodes can use the same hashtable. A node in C code is represented as follows:

```
struct dddnode {
    uint64_t a, b;
} * dddnode_t;
```

In this struct the value(32 bits), the true edge(40 bits), the false edge(40 bits) and a type bit, operator bit and flag bit are stored. These values are not specifically named in the struct, all values are stored in the two integers a and b. Figure 10 shows how this is coded in memory. The type, operator and flag bit are stored in the black areas. We do not show them explicitly due to the scale. The type bit indicates if a node is a DDD or an LDD node, if it is set to 0 it should be treated as a normal LDD node. The operator bit shows if the operator is $<$ or \leq , this can only be used if the type bit is also set to 1(DDD). The flag bit is used in some algorithms to indicate that a certain node has already been visited. All of this is stored compactly in the two 64 bit integers. The total information is 115 bits, so there are still 17 unused bits, all unused bits are set to 0. The depth of the node is not stored, this can be calculated by going down through the structure. This implies that no level can be skipped. Other DDD algorithms and reductions show that some levels are not needed. We solved this by indication a skipped level by $< \infty$, which is true for every upper bound. For such nodes the false edge will always directly lead to the false end node.

4.6 Creating Nodes

To create a node a special MK function is used. This function will ensure that a DDD is always locally reduced. This MK function is shown in algo-




Figure 10: In memory representation of DDD node

rithm 3. This function ensures the correct total structure and puts newly created nodes in the hashtable. The actual creation of a node is done in the MakeNode function that is called inside the MK function. The code for the MakeNode function is not shown here as it is only technical coding, putting all the information in the struct.

Algorithm 3 MK

```

1: procedure MK(value, h, l, type, op)
2:   if  $h = 0 \wedge type = LDD$  then
3:     return l
4:   if  $h = 1 \wedge l = 1$  then
5:     return 1
6:   if  $h = 0 \wedge l = 0$  then
7:     return 0
8:   if  $h = 0 \wedge l \neq 0$  then
9:     return 1
10:  if  $h = high(l)$  then
11:    return l
12:  node = MAKENODE(value, h, l, type, op)
13:  if node  $\notin$  table then
14:    PUT(node)
15:  return node

```


4.7 Minus

The minus function, used for the reachability, has not been implemented as an DDD functions. This function is different to other functions, as information has to be transferred over different levels. For simple cases, an upper-bound in one of the operands of the minus, can become a lower-bound in the result, and vice-versa. A simple one dimensional example is $[0..8]/[0..4]$, this will result in $[4..8]$. In this case the 4 is the upper-bound of the subtrahend. It will however become the lower-bound of the difference. As lower- and upper-bounds are saved on different levels in DDDs this makes the function different from all other functions, which only look at values on the same level.


In figure 11 we have a two-dimensional example of how the minus function can become more complex for multiple-dimensions. In this case we

make a hole in a larger zone. Both the minuend and the subtrahend are represented by a DDD with a single path. For simplicity we removed the diagonals in this example, as they play no role. The difference however becomes a DDD with 4 paths and 10 nodes. Again a lot of upper- and lower-bounds are switched. Already for this example we could not find an algorithm that does this in general. For more dimensions, and DDDs with already multiple paths the problem will only get harder. That is why we returned to a DBM function for this.


The DBM function we use is defined in the Uppaal DBM library. The minus function is defined over a federation of DBMs. This federation is a C++ class containing multiple DBMs. This federation is needed as we can do a minus over a collection of zones, multiple paths in the DDD, and the result can contain multiple zones. As already shown in the example of figure 11. For this function we first take the normal LDD minus function over the discrete part. At the first DDD level representing the zones, the DBM function is called. From this level all possible paths are searched, and for each path a DBM is created. All these DBMs are put in a federation, on which the library function can be called. The result is again (a possibly empty) federation. If the federation is empty, simply a DDD-false node is returned. Otherwise each DBM is turned into a DDD path and these paths are unioned together.



(a) Minuend



(b) Subtrahend



(c) Difference

Figure 11: Minus complexity example




Figure 12: DDD representation of the minuend and subtrahend of figure 11

5 Benchmarks

Below we describe the different models we used to run the benchmarks. We tried to find models that scale up for a number of nodes or processes, so that we can also check the behaviour of our approaches for different sizes of the same model. In this section we use the terminology 'locations' and other 'discrete variables'. The definition of timed automata does not have this difference, but we use it to describe models, because the time extrapolation is dependent on locations, and not on the other discrete variables. This dependency fills a large part of the dependency matrices. In this section, a location is a state in the Uppaal transition system editor, the other discrete variables are declared in the C-like syntax that Uppaal uses.

5.1 Viking

The set of Viking tests models a torch that can be taken from one side to the other by either one or two vikings. By some timing guards, the vikings have to hold one to the torch for a minimal amount of time, given by guards. The model has a low number of discrete variables, one per viking, one for the torch and an indicator for the side on which the torch is. It has a global clock and a clock per viking.

The model results in a densely filled dependency matrix. The torch and all viking variables are always read for the time extrapolation. Only the side indicator is not read always. The write matrix is sparser.

The difference between the LDD and DDD representation is quite small for this model. In the extrapolation step all clock zones are set to $[0..\infty]$ for all states, so in both diagrams the zones are represented by a single path. So the interesting things are only happening in the discrete parts.

5.2 Fischer

Fischer's mutual exclusion protocol [18] is modelled for a number of processes. There is no synchronization between processes, only blocking of actions can occur. This model has a slightly higher number of discrete variables compared to the Viking tests. Each process has a location and an id. The model also has 2 global discrete variables. Each process has a local clock, no global clock is used.

The dependency matrix of this model has some sparse rows, as each model has an id, which is a constant and can only be read. Again all the location variables are always read due to the time extrapolation.

5.3 CSMA-CD

The Carrier Sense Multiple Access/Collision Detection [31] is modelled for a number of senders. The model has a few discrete variables. Each sender

only has a location and only one global counter is used. Each sender and the bus have a local clock. The model uses a lot of synchronizations between the senders and the bus.

5.4 Animo

We could not use the ANIMO models, only the smallest model with no synchronizations was possible. As we started the project to work on ANIMO models, we still included that single model in the benchmark set. It is a model with only one node, so only one location variable. The model has two clocks, a global clock and a clock for the node. Further it does have quite a large number of discrete variables. Both the global declaration and the node have a portion of c-like code with a number of global variables.

This results in a model with a quite sparse dependency matrix, as only the single location is used for the time extrapolation. We expected this model to have good performance for the LDD method with variable re-ordering.

5.5 Lynch-Shavit

The Lynch-Shavit mutual exclusion protocol [20] is modelled for different number of processes. The structure of the model is quite like the Fischer model. It only uses one global variable more than Fischer.

5.6 Milner

Milners scheduler [22] is modelled for a number of nodes. The structure is like that of the CSMA-CD model, except that it does not use a bus. The model has a lot of synchronizations between the nodes, and between the node and a global process. Each node has two clocks, so the zone representation blows up quickly.

5.7 Other models

We also used some models that we could not scale up enough due to memory/time limitations, or that could not scale up due to the nature of the model. We will not describe those models into detail. These models were the critRegion, Critical, bocdp(-fixed) [16], bando and timelock model.


6 Results

There is a difference between the number of nodes for the normal BFS and the BFS without minus. This is possible because we do not use a canonical form of DDDs. Most results show a higher number of nodes for the runs with the minus. In figure 13 we show an example of how this can happen. We assume all zones in the figures belong to the same set of locations. In figure 13a we have the zone that is already visited. Now a new state with the zone in figure 13b is discovered. If the minus is not used, successors of this state are directly generated from the set of locations and this zone. If the minus is used the first zone will first be subtracted before successors are generated. The result of the subtraction is shown in figure 13c. This is not a convex zone, so a DDD with multiple paths is needed. From this state also other successors can be generated, possibly needing more nodes to be represented. If the newly generated states are then unioned with the visited set the result can again have more nodes than the version without minus. The less fractionated zones in the current set can also have implications on the time results, as less work in the next-state function is needed. On the other hand the next-state function can also need extra time, as some states would otherwise have completely been removed from the current set, and no work for that states would need to be done.


In the LDD solution the standard setup with no reordering is for most models faster and uses less nodes than with the reordering algorithms. This is probably due to the densely filled dependency matrices as described in section ???. If we could make these matrices sparser we expect better results from the reordering algorithms.

	DDD		LDD					
	no-min		gsa	rb4w	cw	rs,rn	rs,ru	
fischer1	0.8	0.8	0.8	1.1	0.8	0.8	0.8	
fischer2	0.9	0.9	0.9	1.5	0.9	0.9	0.9	
fischer3	0.9	0.9	0.9	2.0	0.9	0.9	0.9	
fischer4	1.2	1.2	1.2	2.9	1.3	1.2	1.2	
fischer5	14.1	10.2	6.1	9.3	7.3	6.5	5.9	
critRegion1	0.9	0.9	0.9	1.2	0.9	0.9	0.8	
critRegion2	0.9	0.9	0.9	1.5	0.9	0.9	0.9	
critRegion3	3.7	3.6	1.8	2.8	1.9	1.8	2.2	
Critical_01-25-50	0.9	0.9	0.9	1.3	0.9	0.9	0.9	
Critical_02-25-50	0.9	1.0	1.0	1.6	1.0	1.0	1.0	
Critical_03-25-50	13.8	13.8	7.2	9.9	9.3	7.4	6.9	
CSMACD_01	0.8	0.8	0.8	1.0	0.8	0.8	0.8	
CSMACD_02	1.0	1.0	1.0	1.3	1.0	1.0	1.0	
CSMACD_03	1.0	1.0	1.0	1.5	1.0	1.0	1.0	
CSMACD_04	1.2	1.2	1.2	1.7	1.2	1.1	1.2	
CSMACD_05	1.9	1.4	1.5	2.1	1.5	1.5	1.4	
CSMACD_06	5.0	2.0	2.6	3.2	2.6	2.6	2.4	
CSMACD_07	19.7	3.7	7.2	7.6	7.4	7.4	6.5	
CSMACD_08	237.0	10.4	26.7	25.5	28.0	28.0	23.5	
viking1	0.8	0.8	0.8	1.1	0.8	0.8	0.8	
viking2	0.9	0.9	0.9	1.3	0.9	0.9	0.9	
viking3	0.9	0.9	0.9	1.5	0.9	0.9	0.9	
viking4	1.0	1.0	1.0	1.7	1.0	1.0	1.0	
viking5	1.1	1.1	1.1	2.0	1.1	1.1	1.1	
viking6	1.7	1.7	2.0	2.9	2.0	2.0	1.8	
viking7	2.2	2.2	2.5	3.6	2.5	2.5	2.2	
viking8	4.7	4.7	5.8	6.4	5.9	5.9	4.9	
viking9	12.1	12.2	16.2	14.9	16.4	16.3	13.1	
viking10	33.9	34.0	46.9	38.9	47.2	47.3	37.1	
Lynch1-16	0.8	0.8	0.8	1.4	0.8	0.8	0.8	
Lynch2-16	0.9	0.9	0.9	2.2	0.9	0.9	0.9	
Lynch3-16	1.3	1.3	1.2	3.3	1.3	1.2	1.2	
Lynch4-16	8.8	8.6	8.4	11.3	10.0	8.6	8.2	
bocdp	1.6	1.6	1.6	14.6	1.6	1.6	1.6	
bocdpFIXED	1.6	1.6	1.6	13.5	1.6	1.6	1.5	
bando	1.6	1.6	1.5	13.5	1.6	1.6	1.5	
timelock	0.7	0.7	0.7	0.9	0.7	0.7	0.7	
Milner-2Nodes-flat	0.9	0.9	0.9	1.3	0.9	0.9	0.9	
Milner-3Nodes-flat	1.0	0.9	1.0	1.6	1.0	1.0	1.0	
Milner-4Nodes-flat	1.1	1.0	1.6	2.2	1.6	1.6	1.5	
Milner-5Nodes-flat	1.2	1.2	2.1	2.9	2.1	2.1	2.0	
Milner-6Nodes-flat	1.5	1.4	3.0	3.8	3.0	3.0	2.8	
Milner-7Nodes-flat	1.8	1.7	4.2	5.2	4.3	4.3	4.0	
Milner-8Nodes-flat	2.3	2.2	6.1	7.0	6.2	6.2	5.5	
hddi_input.1	0.9	0.9	0.9	1.0	0.9	0.9	0.9	
hddi_input.2	1.0	0.9	0.9	1.2	0.9	0.9	0.9	
hddi_input.3	179.4	1.0	1.1	1.4	1.1	1.1	1.1	
ANIMO_small	1.1	1.1	1.1	1.8	1.1	1.1	1.1	


	DDD		LDD					
	no-minus		gsa		rb4w	cw	rs,rn	rs,ru
fischer1	14	14	14	13	14	13	14	14
fischer2	66	66	66	68	63	66	66	66
fischer3	509	288	409	505	433	532	409	409
fischer4	5025	1300	2541	3905	3190	4184	2541	2541
fischer5	49634	5535	17131	30665	26004	32446	17131	17131
critRegion1	24	24	24	24	20	26	24	24
critRegion2	251	190	243	358	296	242	243	243
critRegion3	4643	3825	3743	5789	5506	4627	3743	3743
Critical.01-25-50	25	25	25	24	23	29	25	25
Critical.02-25-50	313	262	316	500	427	370	316	316
Critical.03-25-50	12322	11183	17505	29297	28443	20293	17505	17505
CSMACD_01	17	17	17	17	17	17	17	17
CSMACD_02	112	108	99	101	101	101	99	99
CSMACD_03	686	458	500	553	551	551	500	500
CSMACD_04	3305	1356	2205	2528	2520	2520	2205	2205
CSMACD_05	13867	3478	8634	10154	10127	10127	8634	8634
CSMACD_06	51633	7925	30862	37022	36938	36938	30862	30862
CSMACD_07	176965	17069	102821	125264	125019	125019	102821	102821
CSMACD_08	569760	36098	324047	329844	398899	398899	324047	324047
vikings1	12	15	15	15	24	24	15	15
vikings2	37	37	37	37	66	66	37	37
vikings3	86	86	86	91	176	176	86	86
vikings4	105	105	105	111	196	196	105	105
vikings5	124	124	124	131	216	216	124	124
vikings6	233	233	233	240	504	504	233	233
vikings7	190	190	190	199	342	342	190	190
vikings8	224	224	224	235	415	415	224	224
vikings9	263	263	263	276	495	495	263	263
vikings10	304	304	304	317	581	581	304	304
Lynch1-16	24	24	24	22	27	21	24	24
Lynch2-16	162	162	173	185	217	187	173	173
Lynch3-16	1175	922	1277	1757	2600	1649	1277	1277
Lynch4-16	14280	8246	11113	22033	32144	17146	11113	11113
bocdp	541	541	572	329	435	517	572	572
bocdpFIXED	542	542	572	428	448	514	572	572
bando	542	542	572	428	448	514	572	572
timelock	4	4	4	4	4	4	4	4
Milner-2Nodes-flat	442	245	338	327	394	338	338	338
Milner-3Nodes-flat	2709	918	1602	1571	1702	1602	1602	1602
Milner-4Nodes-flat	4999	2968	4834	4789	4997	4834	4834	4834
Milner-5Nodes-flat	9106	5293	8653	8596	8946	8653	8653	8653
Milner-6Nodes-flat	17008	7755	14100	14018	14551	14100	14100	14100
Milner-7Nodes-flat	25493	12188	21455	21343	22098	21455	21455	21455
Milner-8Nodes-flat	39887	16324	31008	30884	31874	31008	31008	31008
hddi.input_1	221	119	130	134	134	134	130	130
hddi.input_2	2735	693	1021	1025	1090	1023	1021	1021
hddi.input_3	20485	2013	3675	3675	3971	3675	3675	3675
ANIMO_small	235	235	197	191	405	185	197	197



(a) Visited Zone



(b) Current Zone



(c) After Minus

Figure 13: Minus fragmentation

References

- [1] S. B. Akers. Binary decision diagrams. *IEEE Trans. Comput.*, 27(6):509–516, June 1978.
- [2] Rajeev Alur and David L. Dill. A theory of timed automata. *Theoretical Computer Science*, 126(2):183 – 235, 1994.
- [3] Gerd Behrmann, Patricia Bouyer, Kim G. Larsen, and Radek Pelánek. *Lower and Upper Bounds in Zone Based Abstractions of Timed Automata*, pages 312–326. Springer Berlin Heidelberg, Berlin, Heidelberg, 2004.
- [4] Gerd Behrmann, Alexandre David, and Kim G. Larsen. A tutorial on Uppaal. In Marco Bernardo and Flavio Corradini, editors, *Formal Methods for the Design of Real-Time Systems*, volume 3185 of *Lecture Notes in Computer Science*, pages 200–236. Springer Berlin Heidelberg, 2004.
- [5] Johan Bengtsson. *Clocks, DBMS and States in Timed Systems (Uppsala Dissertations from the Faculty of Science Technology, 39)*. Uppsala Universitet, 7 2002.
- [6] Dirk Beyer. Efficient reachability analysis and refinement checking of timed automata using BDDs. In T. Margaria and T. F. Melham, editors, *Proceedings of the 11th IFIP Advanced Research Working Conference on Correct Hardware Design and Verification Methods (CHARME 2001, Livingston, September 4-7)*, LNCS 2144, pages 86–91. Springer-Verlag, Heidelberg, 2001.
- [7] Dirk Beyer, Claus Lewerentz, and Andreas Noack. Rabbit: A tool for BDD-based verification of real-time systems. In W. A. Hunt and F. Somenzi, editors, *Proceedings of the 15th International Conference on Computer Aided Verification (CAV 2003, Boulder, CO, July 8-12)*, LNCS 2725, pages 122–125. Springer-Verlag, Heidelberg, 2003.
- [8] S. C. C. Blom, J. C. van de Pol, and M. Weber. LTSmin: Distributed and symbolic reachability. In T. Touili, B. Cook, and P. Jackson, editors, *Computer Aided Verification, Edinburgh*, volume 6174 of *Lecture Notes in Computer Science*, pages 354–359, Berlin, July 2010. Springer Verlag.
- [9] Stefan Blom and Jaco van de Pol. Symbolic reachability for process algebras with recursive data types. In J.S. Fitzgerald, A.E. Haxthausen, and H. Yenigun, editors, *Theoretical Aspects of Computing*, volume 5160 of *Lecture Notes in Computer Science*, pages 81–95, Berlin, Germany, August 2008. Springer Verlag.

- [10] R.E. Bryant. Graph-based algorithms for boolean function manipulation. *Computers, IEEE Transactions on*, C-35(8):677–691, Aug 1986.
- [11] A. E. Dalsgaard, A. W. Laarman, K. G. Larsen, M. C. Olesen, and J. C. van de Pol. Multi-core reachability for timed automata. In M. Jurdzinski and D. Nickovic, editors, *10th International Conference on Formal Modeling and Analysis of Timed Systems, FORMATS 2012, London, UK*, volume 7595 of *Lecture Notes in Computer Science*, pages 91–106, London, September 2012. Springer Verlag.
- [12] Andreas Engelbrecht Dalsgaard, Ren Rydhof Hansen, Kenneth Yrke Jørgensen, Kim Gulstrand Larsen, Mads Chr. Olesen, Petur Olsen, and Ji Srba. opaal: A lattice model checker. In Mihaela Bobaru, Klaus Havelund, GerardJ. Holzmann, and Rajeev Joshi, editors, *NASA Formal Methods*, volume 6617 of *Lecture Notes in Computer Science*, pages 487–493. Springer Berlin Heidelberg, 2011.
- [13] David L. Dill. Timing assumptions and verification of finite-state concurrent systems. In Joseph Sifakis, editor, *Automatic Verification Methods for Finite State Systems*, volume 407 of *Lecture Notes in Computer Science*, pages 197–212. Springer Berlin Heidelberg, 1990.
- [14] Junwei Du, Huiping Zhang, Gang Yu, and Xi Wang. A full symbolic compositional reachability analysis of timed automata based on BDD. In *Advanced Computational Intelligence (ICACI), 2015 Seventh International Conference on*, pages 218–222, March 2015.
- [15] R. Ehlers, D. Fass, M. Gerke, and H.-J. Peter. Fully symbolic timed model checking using constraint matrix diagrams. In *Real-Time Systems Symposium (RTSS), 2010 IEEE 31st*, pages 360–371, Nov 2010.
- [16] K. Havelund, A. Skou, K. G. Larsen, and K. Lund. Formal modeling and analysis of an audio/video protocol: an industrial case study using UPPAAL. In *Real-Time Systems Symposium, 1997. Proceedings., The 18th IEEE*, pages 2–13, Dec 1997.
- [17] A. W. Laarman, J. C. van de Pol, and M. Weber. Multi-Core LTSmin: Marrying Modularity and Scalability. In M. Bobaru, K. Havelund, G. Holzmann, and R. Joshi, editors, *Proceedings of the Third International Symposium on NASA Formal Methods, NFM 2011, Pasadena, CA, USA*, volume 6617 of *Lecture Notes in Computer Science*, pages 506–511, Berlin, July 2011. Springer Verlag.
- [18] Leslie Lamport. A fast mutual exclusion algorithm. *ACM Trans. Comput. Syst.*, 5(1):1–11, January 1987.
- [19] Kim Larsen, Carsten Weise, Wang Yi, and Justin Pearson. Clock difference diagrams. *BRICS Report Series*, 5(46), 1998.

- [20] N. Lynch and N. Shavit. Timing based mutual exclusion. In *Proc. of the Annual Real-Time Symposium (RTSS)*, pages 2–11, 1992.
- [21] Jeroen Meijer, Gijs Kant, Stefan Blom, and Jaco van de Pol. Read, write and copy dependencies for symbolic model checking. In Eran Yahav, editor, *Hardware and Software: Verification and Testing*, volume 8855 of *Lecture Notes in Computer Science*, pages 204–219. Springer International Publishing, 2014.
- [22] R. Milner. *Communication and Concurrency*. Prentice-Hall, Inc., Upper Saddle River, NJ, USA, 1989.
- [23] J. Møller, J. Lichtenberg, H. R. Andersen, and H. Hulgaard. Difference decision diagrams. Technical Report IT-TR-1999-023, Department of Information Technology, Technical University of Denmark, Building 344, DK-2800 Lyngby, Denmark, February 1999.
- [24] Jesper Møller, Henrik Hulgaard, and Henrik Reif Andersen. Symbolic model checking of timed guarded commands using difference decision diagrams. *The Journal of Logic and Algebraic Programming*, 5253:53 – 77, 2002.
- [25] Jesper Møller, Jakob Lichtenberg, HenrikReif Andersen, and Henrik Hulgaard. Difference decision diagrams. In Jörg Flum and Mario Rodriguez-Artalejo, editors, *Computer Science Logic*, volume 1683 of *Lecture Notes in Computer Science*, pages 111–125. Springer Berlin Heidelberg, 1999.
- [26] TruongKhanh Nguyen, Jun Sun, Yang Liu, JinSong Dong, and Yan Liu. Improved BDD-based discrete analysis of timed systems. In Dimitra Giannakopoulou and Dominique Méry, editors, *FM 2012: Formal Methods*, volume 7436 of *Lecture Notes in Computer Science*, pages 326–340. Springer Berlin Heidelberg, 2012.
- [27] Stefano Schivo, Jetse Scholma, Brend Wanders, Ricardo A. Urquidi Camacho, Paul E. van der Vet, Marcel Karperien, Rom Langerak, Jaco van de Pol, and Janine N. Post. Modelling biological pathway dynamics with timed automata. In *12th IEEE International Conference on Bioinformatics & Bioengineering, BIBE 2012, Larnaca, Cyprus, November 11-13, 2012*, pages 447–453, 2012.
- [28] A. Srinivasan, T. Ham, S. Malik, and R.K. Brayton. Algorithms for discrete function manipulation. In *Computer-Aided Design, 1990. ICCAD-90. Digest of Technical Papers., 1990 IEEE International Conference on*, pages 92–95, Nov 1990.

- [29] Tom van Dijk and Jaco van de Pol. Sylvan: Multi-core decision diagrams. In Christel Baier and Cesare Tinelli, editors, *Tools and Algorithms for the Construction and Analysis of Systems*, volume 9035 of *Lecture Notes in Computer Science*, pages 677–691. Springer Berlin Heidelberg, 2015.
- [30] Farn Wang. Efficient verification of timed automata with BDD-like data-structures. In LenoreD. Zuck, PaulC. Attie, Agostino Cortesi, and Supratik Mukhopadhyay, editors, *Verification, Model Checking, and Abstract Interpretation*, volume 2575 of *Lecture Notes in Computer Science*, pages 189–205. Springer Berlin Heidelberg, 2003.
- [31] Sergio Yovine. Kronos: a verification tool for real-time systems. *International Journal on Software Tools for Technology Transfer*, 1(1-2):123–133, 1997.
- [32] Huiping Zhang, Junwei Du, Ling Cao, and Guixin Zhu. A full symbolic reachability analysis algorithm of timed automata based on BDD. In *Autonomous Decentralized Systems (ISADS), 2015 IEEE Twelfth International Symposium on*, pages 301–304, March 2015.